



Harshit Rai, Ishank Vasania, Mohammad Danish Nadeem

# Main Goal

- Demonstrate various features of MongoDB
- Test the effectiveness of MongoDB in handling graph data

# Features of MongoDB

- Flexible Schema
- Scalability
- High Performance
- Rich Query Language
- Replication and High Availability
- Horizontal Scaling
- Easy Integration
- Full Text Search
- Geospatial Capabilities
- Aggregation Framework
- Document-Oriented

# Applications of MongoDB

1. Real-time Analytics
2. Internet of Things (IoT)
3. Catalog and Product Inventories
4. Social Networks and Collaboration
5. Data Caching and Session Storage
6. Mobile Application

Famous Companies using MongoDB: Ebay, EA, Shutterfly and many others

# MongoDB Compass

## Advantages:

- User-Friendly Interface
- Schema Visualization
- Query Building
- Index Optimization

## Disadvantages:

- Limited Functionality
- Learning Curve

Documents  
data-514.amazon...

My Queries

Databases



Search

data-514

amazon-meta

## data-514.amazon-meta

548.6k 0  
DOCUMENTS INDEXES

Documents

Aggregations

Schema

Explain Plan

Indexes

Validation

Filter



Type a query: { field: 'value' }

Reset

Find



More Options

ADD DATA

EXPORT COLLECTION

1 - 20 of 548552



```
_id: ObjectId('64682d17d2cff919f431f886')
Id: 0
ASIN: "0771044445"
```

```
_id: ObjectId('64682d17d2cff919f431f887')
Id: 1
ASIN: "0827229534"
title: "Patterns of Preaching: A Sermon Sampler"
group: "Book"
salesrank: 396585
similar: Array
categories: Array
reviews: Object
```

```
_id: ObjectId('64682d17d2cff919f431f888')
Id: 2
ASIN: "0738700797"
title: "Candlemas: Feast of Flames"
group: "Book"
salesrank: 168596
similar: Array
categories: Array
reviews: Object
```

```
_id: ObjectId('64682d17d2cff919f431f889')
Id: 3
ASIN: "0486287785"
title: "World War II Allied Fighter Planes Trading Cards"
group: "Book"
salesrank: 1270652
similar: Array
categories: Array
reviews: Object
```

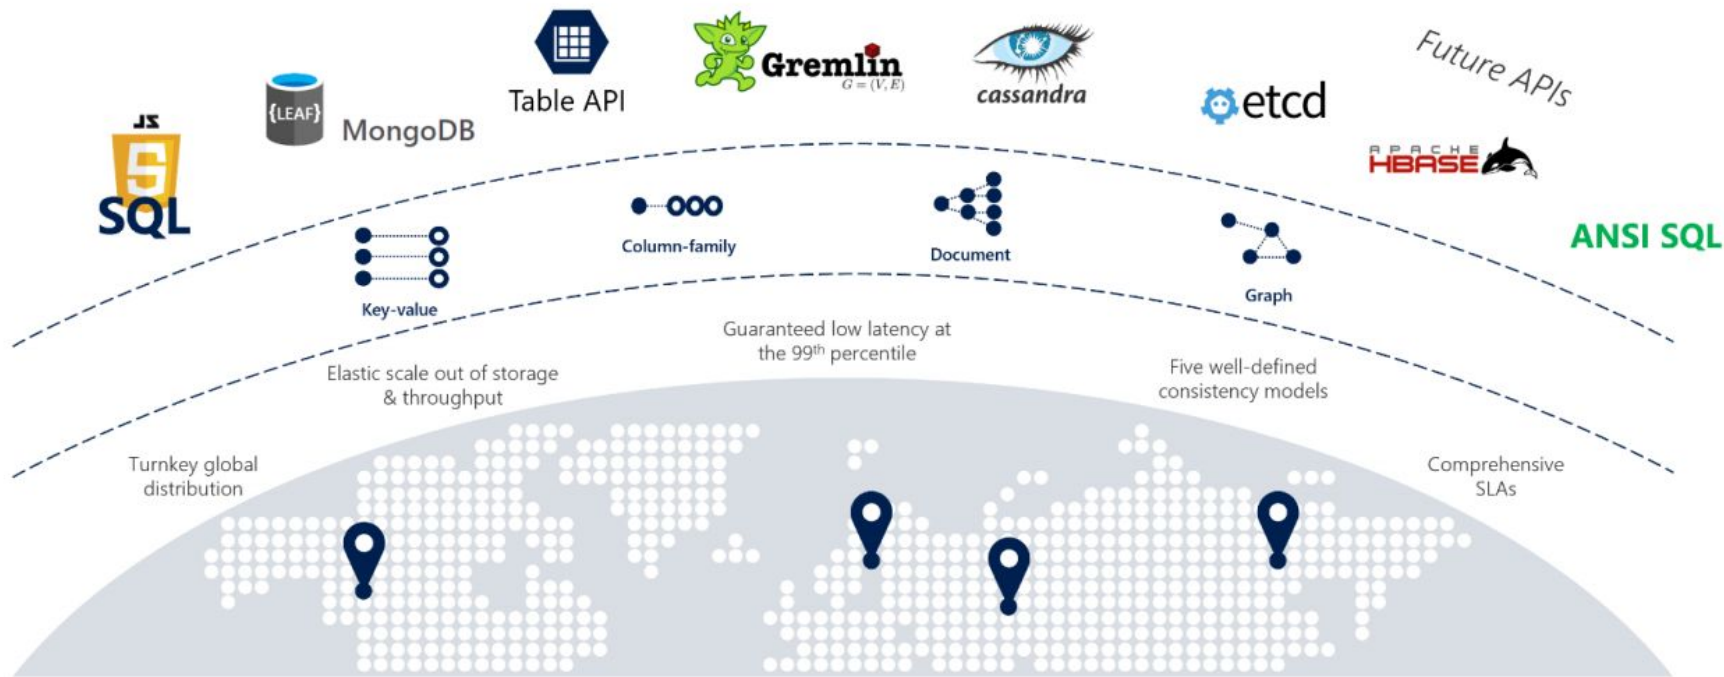
```
_id: ObjectId('64682d17d2cff919f431f88a')
```

# MongoDB Atlas v/s Cosmos DB

| Comparison Criteria  | Cosmos DB                                 | Atlas   |
|----------------------|---|---|
| Vendor and Ecosystem | Microsoft Azure                           | MongoDB   |
| Supported Databases  | DocumentDB, MongoDB, SQL, Gremlin, Table  | MongoDB   |
| Compatibility        | Compatible with multiple APIs             | Compatible with MongoDB's wire protocol           |
| Global Distribution  | Multiple data centers, global replication | Multiple regions, data replication across regions |
| Pricing Model        | Provisioned throughput                    | Usage-based pricing                               |

# Azure Cosmos DB


Microsoft's globally distributed, massively scalable, multi-model database service





# Sharding in MongoDB

- Sharding: MongoDB's method for distributing data across multiple machines.
- Challenges: Large data sets or high throughput can strain a single server's capacity.
- Vertical Scaling: Increasing a server's capacity (CPU, RAM, storage), but with limitations.
- **Horizontal Scaling**: Dividing workload among multiple servers, providing efficiency and flexibility.
- Trade-off: Horizontal scaling requires more complex infrastructure but can be cost-effective.
- Sharding != Partitioning



# ACID in MongoDB

1. Atomicity, Consistency, Isolation, and Durability.
2. Initial support for ACID principles on single-document level.
3. Added support for multi-document ACID transaction in v4.0 (2018).
4. Added support for distributed multi-document ACID transaction in v4.2 (2019).
5. Application of multi-document ACID transactions-
  - a. Updating the status of an account across all those users' documents.
  - b. Logging custom application actions -- say when a user transfers ownership of an entity, the write should not be successful if the logging isn't.
  - c. Aggregations of hundreds of thousands of trades, need to be updated every time trades are added or modified.

# Data Profile

**Amazon-metadata** file with about 550000 rows containing information on the products sold on Amazon. The size of the file is ~1 GB.

## Fields and their descriptions:

- Id: A unique identifier for the product, assigned by Amazon.
- ASIN: The Amazon Standard Identification Number for the product.
- Title: The title of the product.
- Group: The group that the product belongs to (e.g., Book, Music, DVD).
- Salesrank: The sales rank of the product on Amazon.
- **Similar**: A list of ASINs for products that are frequently co-purchased with the current product.
- Categories: A list of category paths that the product belongs to on Amazon.
- Reviews: A list of review objects, each containing information about a customer review for the product.

```

Id: 1
ASIN: 0827229534
  title: Patterns of Preaching: A Sermon Sampler
  group: Book
  salesrank: 396585
  similar: 5 0804215715 156101074X 0687023955 0687074231 082721619X
  categories: 2
    |Books[283155]|Subjects[1000]|Religion & Spirituality[22]|Christianity[12290]|Clergy[12360]|Preaching[12368]
    |Books[283155]|Subjects[1000]|Religion & Spirituality[22]|Christianity[12290]|Clergy[12360]|Sermons[12370]
  reviews: total: 2 downloaded: 2 avg rating: 5
    2000-7-28  cutomer: A2JW670Y8U6HHK  rating: 5  votes: 10  helpful: 9
    2003-12-14  cutomer: A2VE83MZ98ITY  rating: 5  votes: 6  helpful: 5

```

```

Id: 2
ASIN: 0738700797
  title: Candlemas: Feast of Flames
  group: Book
  salesrank: 168596
  similar: 5 0738700827 1567184960 1567182836 0738700525 0738700940
  categories: 2
    |Books[283155]|Subjects[1000]|Religion & Spirituality[22]|Earth-Based Religions[12472]|Wicca[12484]
    |Books[283155]|Subjects[1000]|Religion & Spirituality[22]|Earth-Based Religions[12472]|Witchcraft[12486]
  reviews: total: 12 downloaded: 12 avg rating: 4.5
    2001-12-16  cutomer: A11NC06YTE4BTJ  rating: 5  votes: 5  helpful: 4
    2002-1-7   cutomer: A9CQ3PLRNIR83  rating: 4  votes: 5  helpful: 5
    2002-1-24  cutomer: A13SG9ACZ905IM  rating: 5  votes: 8  helpful: 8
    2002-1-28  cutomer: A18DAI6VEYMAZA  rating: 5  votes: 4  helpful: 4
    2002-2-6   cutomer: A2P6KAWXJ16234  rating: 4  votes: 16  helpful: 16
    2002-2-14  cutomer: AMACWC3M7PQFR  rating: 4  votes: 5  helpful: 5
    2002-3-23  cutomer: A3G07UV9XX14D8  rating: 4  votes: 6  helpful: 6
    2002-5-23  cutomer: A1GIL64QK68WKL  rating: 5  votes: 8  helpful: 8
    2003-2-25  cutomer: AE0B0F20NQJWV  rating: 5  votes: 8  helpful: 5
    2003-11-25 cutomer: A3IGHTES8ME05L  rating: 5  votes: 5  helpful: 5
    2004-2-11  cutomer: A1CP26N8RHYVVO  rating: 1  votes: 13  helpful: 9
    2005-2-7   cutomer: ANEIANH0WAT9D  rating: 5  votes: 1  helpful: 1

```

```

Id: 3
ASIN: 0486287785
  title: World War II Allied Fighter Planes Trading Cards
  group: Book
  salesrank: 1270652
  similar: 0
  categories: 1
    |Books[283155]|Subjects[1000]|Home & Garden[48]|Crafts & Hobbies[5126]|General[5144]

```

```

_id: ObjectId('64682d17d2cff919f431f887')
Id: 1
ASIN: "0827229534"
  title: "Patterns of Preaching: A Sermon Sampler"
  group: "Book"
  salesrank: 396585
  similar: Array
    0: "0804215715"
    1: "156101074X"
    2: "0687023955"
    3: "0687074231"
    4: "082721619X"
  categories: Array
    0: Array
      0: "Books[283155]"
      1: "Subjects[1000]"
      2: "Religion & Spirituality[22]"
      3: "Christianity[12290]"
      4: "Clergy[12360]"
      5: "Preaching[12368]"
    1: Array
      0: "Books[283155]"
      1: "Subjects[1000]"
      2: "Religion & Spirituality[22]"
      3: "Christianity[12290]"
      4: "Clergy[12360]"
      5: "Sermons[12370]"
  reviews: Object
    total: 2
    downloaded: 2
    avg rating: "5"
  customer_reviews: Array
    0: Object
      date: "2000-7-28"
      cutomer: "A2JW670Y8U6HHK"
      rating: 5
      votes: 10
      helpful: 9
    1: Object
      date: "2003-12-14"
      cutomer: "A2VE83MZ98ITY"
      rating: 5
      votes: 6
      helpful: 5

```

Raw text file

Processed JSON file

4 **node** files which represent different combinations of products bought together.

**Node:** Each node in the network represents a unique product. Each node has associated attributes such as product ID, product name, category, brand, description, etc.

**Edge:** The edges in the network represent co-purchasing relationships between products.

An edge exists between two nodes if those products have been purchased together by Amazon customers.

#### Fields and their descriptions:

- Date: The date that the review was posted.
- Customer: A unique identifier for the customer who posted the review.
- Rating: The rating that the customer gave the product (1-5 stars).
- Votes: The number of helpful votes that the review received.
- Helpful: The number of customers who found the review helpful.

```

{
  "From": 0,
  "To": 1
},
{
  "From": 0,
  "To": 2
},
{
  "From": 0,
  "To": 3
},
{
  "From": 0,
  "To": 4
},
{
  "From": 0,
  "To": 5
},
{
  "From": 0,
  "To": 6
},
{
  "From": 0,
  "To": 7
}

```

Raw JSON file

```

{
  "product_id": 0,
  "co_purchases": [
    1,
    2,
    3,
    4,
    5
  ],
  "Date": "March 02 2003",
  "month": "March"
},
{
  "product_id": 1,
  "co_purchases": [
    0,
    2,
    4,
    5,
    15
  ],
  "Date": "March 02 2003",
  "month": "March"
},
{
  "product_id": 2,
  "co_purchases": [
    0,
    11,
    12,

```

Processed JSON file

# Data Preprocessing

**Objective:** Convert the raw data into a semi-structured and usable JSON format.

- Format of the raw data: tab-spaced .txt file
- This format is not supported by MongoDB
- Expected structure is JSON
- JSON format allows for easy data manipulation, querying, and integration with various tools and technologies
- Solution- A custom Python script that processes each product from the raw text file and converts it into a corresponding JSON object.

# Preprocessing Steps

## Amazon-Metadata

### 1. Data Extraction

- a. Read the raw text file containing the Amazon-Meta dataset.
- b. Load the entire dataset into memory.

### 2. Regular Expression Matching

- a. Apply regular expressions to extract relevant information from the dataset.
- b. Use different regex patterns to match global block, review, and customer review patterns.

### 3. JSON Conversion

- a. Parse the extracted information and convert it into a structured JSON format.
- b. Create nested JSON objects and arrays to represent different data components.

### 4. Write to file



## Final JSON Structure:

- Top-level object: Contains an array of individual product objects.
- Product object: Represents a single product with its attributes and reviews.
- Attributes: ASIN, title, group, sales rank, similar products, and categories.
- Similar: ASIN of similar products.
- Categories: Categories to which the product was mapped
- Reviews: Includes total reviews, downloaded reviews, average rating, and customer reviews.

```
_id: ObjectId('64682d17d2cff919f431f887')
Id: 1
ASIN: "0827229534"
title: "Patterns of Preaching: A Sermon Sampler"
group: "Book"
salesrank: 396585
▼ similar: Array
  0: "0804215715"
  1: "156101074X"
  2: "0687023955"
  3: "0687074231"
  4: "082721619X"
▼ categories: Array
  ▼ 0: Array
    0: "Books[283155]"
    1: "Subjects[1000]"
    2: "Religion & Spirituality[22]"
    3: "Christianity[12290]"
    4: "Clergy[12360]"
    5: "Preaching[12368]"
  ▼ 1: Array
    0: "Books[283155]"
    1: "Subjects[1000]"
    2: "Religion & Spirituality[22]"
    3: "Christianity[12290]"
    4: "Clergy[12360]"
    5: "Sermons[12370]"
▼ reviews: Object
  total: 2
  downloaded: 2
  avg rating: "5"
▼ customer_reviews: Array
  ▼ 0: Object
    date: "2000-7-28"
    customer: "A2JW670Y8U6HHK"
    rating: 5
    votes: 10
    helpful: 9
  ▼ 1: Object
    date: "2003-12-14"
    customer: "A2VE83MZ98ITY"
    rating: 5
    votes: 6
    helpful: 5
```

## Node Files

### 1. Load JSON data

- a. Extract the month from the file name
- b. Create a product dictionary
- c. Iterate over each line in the data
- d. Extract "From" and "To" IDs
- e. Add the "To" ID to the co-purchases list of the "From" ID in the dictionary

### 2. Transform and store JSON:

- a. Create a list of products
- b. For each product in the dictionary:
- c. Assign product ID, co-purchases, date, and month
- d. Append the product to the list

**Output:** co-purchases.json that represents co-purchase relationships between products that is useful for further analysis and processing

# Mongodb System Summary

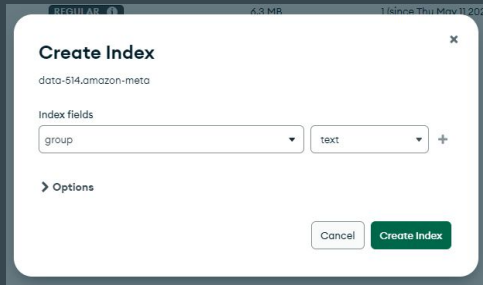
## Data Model

MongoDB uses a document-based data model. Documents are stored as JSON-like objects, and each document can have a different structure. Each document contains key-value pairs, where the keys represent field names and the values represent data. MongoDB supports dynamic schema, which means that we can add new fields to a document at any time. Below is an example of how data can be stored in a MongoDB document-based data *collection*. The Data Model used for the Amazon dataset is explained in more detail in the later part of this document

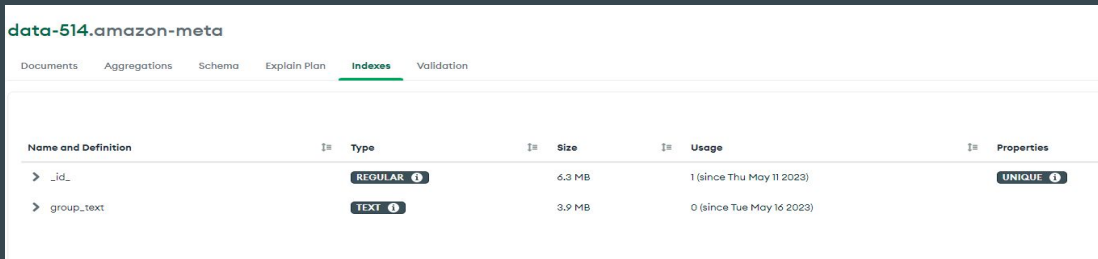
# MongoDB System Summary

## Indexes

MongoDB supports various types of indexes, including single field indexes, compound indexes, multikey indexes, geospatial indexes, and text indexes. Indexes can significantly improve query performance by reducing the number of documents that need to be scanned. Below is an example of an index created on the *group* field of the dataset



The image shows a 'Create Index' dialog box in MongoDB. The database name is 'data-514.amazon-meta'. Under 'Index fields', the field 'group' is selected with a dropdown menu, and the index type is set to 'text' with another dropdown menu. There is a plus sign to add more fields. Below this, there is an 'Options' section with a right-pointing arrow. At the bottom, there are 'Cancel' and 'Create Index' buttons.



The image shows the 'Indexes' tab in the MongoDB interface for the 'data-514.amazon-meta' database. The tabs at the top are 'Documents', 'Aggregations', 'Schema', 'Explain Plan', 'Indexes' (which is active), and 'Validation'. Below the tabs is a table listing the indexes.

| Name and Definition | Type      | Size   | Usage                     | Properties |
|---------------------|-----------|--------|---------------------------|------------|
| > _id_              | REGULAR ⓘ | 6.3 MB | 1 (since Thu May 11 2023) | UNIQUE ⓘ   |
| > group_text        | TEXT ⓘ    | 3.9 MB | 0 (since Tue May 16 2023) |            |

# System Design

## Storing Data in MongoDB

The following steps will be followed to store the dataset in MongoDB.

1. **Initialize System:** Setup the MongoDB system on a local machine and create a corresponding MongoDB instance on Microsoft Azure.
2. **Data Overview:**
  - a. Exploring the Raw Data
  - b. Identifying Useful Fields
  - c. Designing the Data Model
3. **Data Preprocessing and Conversion:** Perform preprocessing to convert the raw data in text format to JSON format using Python. This will help loading the data into a format supported by MongoDB and the data model designed in the previous step.
4. **Data Ingestion:** Create a MongoDB database and collection to store the dataset. We use the MongoDB (Compass) graphical user interface to create the database and collection.
5. **Querying the Data:** Once the data is stored in MongoDB, perform various queries and analysis on the data using the supported query language and aggregation framework.

# MongoDB Data Model Variants

- Embedded Data Model
  - In this variant, related data is stored within a single document as nested fields or arrays.
  - It allows for efficient retrieval of complete data in a single query.
  - Suitable for one-to-one and one-to-many relationships where the embedded data is small and does not require frequent updates.
- Referenced Data Model
  - In this variant, data is distributed across multiple collections, and relationships are established using references or foreign keys.
  - It allows for flexibility in managing and updating related data independently.
  - Suitable for many-to-many relationships or scenarios where the related data is large, frequently updated, or subject to changes.
- Hybrid Data Model
  - This variant combines elements of both the embedded and referenced data models.
  - It leverages the benefits of both approaches, allowing for flexibility and performance optimizations.
  - Suitable for complex relationships or scenarios where a balance between data duplication and data consistency is required.

# Data Model Variants

## Variant 1

Edges collection:

```
{
  "_id": ObjectId("61a76f104e5f151a482f9844"),
  "product_id":
ObjectId("61a76f104e5f151a482f9845"),
  "co_purchases": [
    ObjectId("61a76f104e5f151a482f9846"),
    ObjectId("61a76f104e5f151a482f9847")
  ],
  "date": ISODate("2023-05-20T00:00:00Z"),
  "month": "May"
}
```

## Selected Model

Edges collection:

```
{
  "_id": "ObjectId",
  "product_id": "NumberInt",
  "co_purchases": [
    "NumberInt"
  ],
  "date": "ISODate",
  "month": "String"
}
```

# Data Variants

## Variant 2:

```
{
  "_id":
  ObjectId("61a76f104e5f151a482f
  9844"),
  "From": NumberInt(12345),
  "To": NumberInt(12345)
}
```

## Selected Model

Edges collection:

```
{
  "_id": "ObjectId",
  "product_id": "NumberInt",
  "co_purchases": [
    "NumberInt"
  ],
  "date": "ISODate",
  "month": "String"
}
```



## Variant 2:

dbms

| March 02 2003 |    |       |   |       |
|---------------|----|-------|---|-------|
| id            | pk | old   | * | (1,1) |
| From          |    | int32 | * |       |
| To            |    | int32 | * |       |

| amazon-metadata  |    |       |   |       |
|------------------|----|-------|---|-------|
| id               | pk | old   | * | (1,1) |
| id               |    | int32 | * | (1,1) |
| ASIN             |    | str   | * | (1,1) |
| categories       |    | arr   |   |       |
| [0]              |    | arr   |   |       |
| [0]              |    | str   |   |       |
| group            |    | str   |   |       |
| reviews          |    | doc   |   |       |
| avg rating       |    | str   | * |       |
| customer_reviews |    | arr   | * |       |
| [0]              |    | doc   |   |       |
| customer         |    | str   | * |       |
| date             |    | date  | * |       |
| helpful          |    | int32 | * |       |
| rating           |    | int32 | * |       |
| votes            |    | int32 | * |       |
| downloaded       |    | int32 | * |       |
| total            |    | int32 | * |       |
| salesrank        |    | int32 | * |       |
| similar          |    | arr   |   |       |
| [0]              |    | str   |   |       |
| title            |    | str   |   |       |

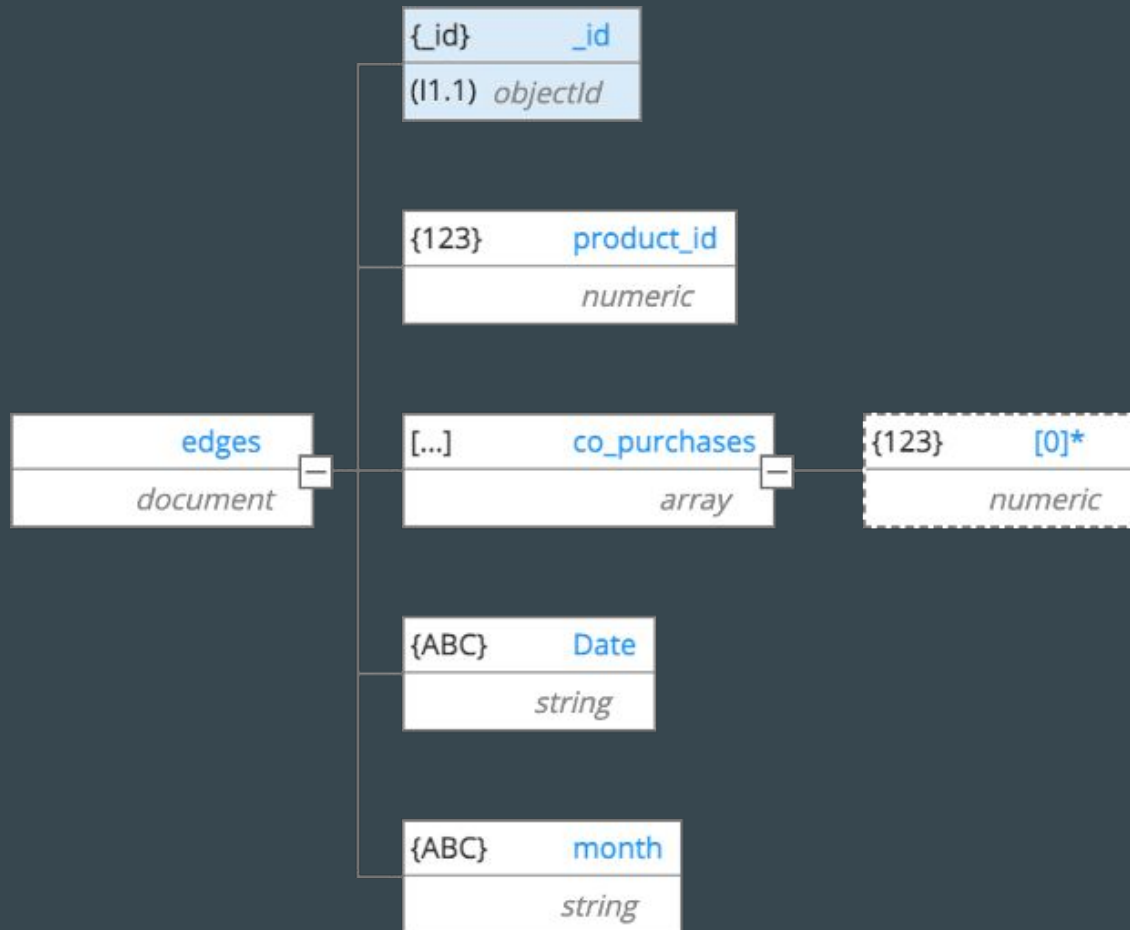
# System Design

The diagram illustrate the schema of the *amazon-metadata* collection that are part of the Amazon dataset under consideration.



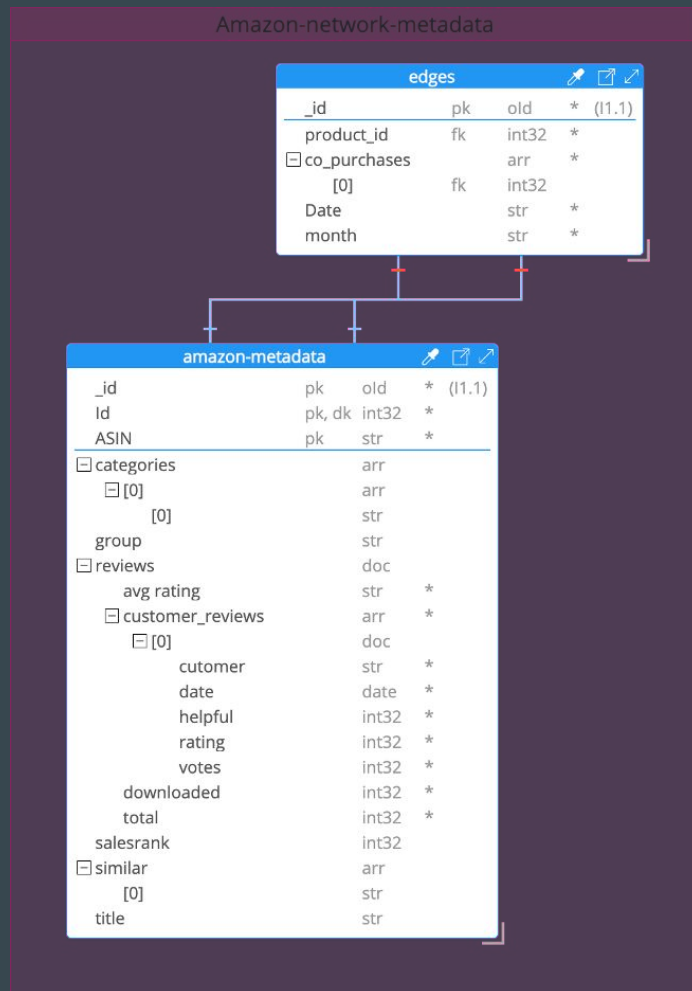
# System Design

The diagram illustrates the schema of the *edges* collection that are part of the Amazon dataset under consideration.



# System Design

## ER/Class Diagram



# System Design

## Indexes used

### Amazon-metadata collection

1. Unique Index on ASIN: The **ASIN** field is unique for each document in the **products** collection, hence we created a unique index on this field to prevent duplicate products from being inserted.
2. Unique Index on Product ID: The **product\_id** field is unique for each document in the **products** collection, hence we created a unique index on this field to prevent duplicate products from being inserted.

# System Design

## Reason for selecting the data model and its potential variant

While designing the data model, the key considerations include the performance requirements of the application, the complexity of the data relationships, and the need for data consistency and integrity.

The original dataset had 4 different files to represent *edges*. If we included these 4 files directly into our mongodb database the size limit was exceeded as a total of 5 collections were found to be too intensive for MongoDB to handle given the resource constraints. In the variant, as could be seen in the variant 2 image, we had the edges from 02-March-2003. To tackle this issue, all of the 4 files containing *edges* collection were merged into one file. The new resultant file created is shown in the ER diagrams in the previous parts and now also has date as well as co-purchases. We basically converted the data from long format to wide format hence saving memory consumption.

In conclusion, for the Amazon product review dataset, we chose a document-based data model in MongoDB due to its flexibility, scalability, and ability to handle unstructured data efficiently.

# System Design

## Preprocessing to restructure the data

The "amazon-meta.txt" input file contains data that our custom Python script analyses and converts to JSON format. The input file includes details on different items sold on Amazon's website. To extract particular patterns, including customer reviews, overall reviews, and information about global blocks, our script makes use of regular expressions. The aforementioned Python script handles a range of keys and values, transforms some values into integers or floats, and arranges the information into dictionaries. The parsed data is then written in JSON format to a new file called "sample.json" and kept in a list. Overall, the script organizes and extracts the pertinent data from the input file into a more understandable and consistent format for additional processing or analysis.

# Aggregations in MongoDB

In MongoDB, aggregation and querying are two fundamental operations used to retrieve and manipulate data.

## Querying:

Querying in MongoDB involves searching for documents in a collection that match certain criteria. It allows you to specify conditions, projections, and sorting parameters to filter and retrieve data. MongoDB provides a flexible and powerful query language that supports a wide range of operators and expressions for complex querying.

```
db.users.find({ name: "John" })
```

## Aggregation:

Aggregation in MongoDB involves processing and transforming data from multiple documents in a collection to produce aggregated results. It allows you to perform operations such as grouping, filtering, sorting, calculating averages, sums, counts, and other aggregations on the data.


```
db.users.aggregate([  
  
  { $group: { _id: "$city", averageAge: { $avg: "$age" } } }  
  
])
```

## Query to answer:

What are the percentages of each rating digit for the product with id: 21?



# Aggregations and Querying



 **CosmoDb** ...

Aggregations  
Amazon-network...


+  
548.6k 0  
DOCUMENTS INDEXES



Amazon-network-metadata.amazon-metadata

Documents Aggregations Schema Explain Plan Indexes Validation

Pipeline  \$match \$unwind \$project  Edit

Explain Export Run More Options ▶

ALL RESULTS  OUTPUT OPTIONS

Showing 1 – 20 [count results](#) < > VIEW  

```
{
  "_id": {},
  "Id": 21,
  "ASIN": "0790747324",
  "title": "The Time Machine",
  "rating_percentage": 100
}
```

```
{
  "_id": {},
  "Id": 21,
  "ASIN": "0790747324",
  "title": "The Time Machine",
  "rating_percentage": 80
}
```

```
{
  "_id": {},
  "Id": 21,
```



CosmoDb

...

Aggregations  
Amazon-network...

My Queries

Databases



Search

Amazon-network-metadata

amazon-metadata

edges

## Amazon-network-metadata.amazon-metadata

548.6k

0

DOCUMENTS

INDEXES

Documents

Aggregations

Schema

Explain Plan

Indexes

Validation

Pipeline



\$match

\$unwind

\$project

Explain

Export

Run

More Options

Problem 1 - modified

SAVE

CREATE NEW

EXPORT TO LANGUAGE

PREVIEW

STAGES

TEXT



```
1 [
2   {
3     $match:
4     {
5       Id: 21,
6     },
7   },
8   {
9     $unwind:
10    "$reviews.customer_reviews",
11  },
12  {
13    $project:
14    {
15      Id: 1,
16      ASIN: 1,
17      title: 1,
18      rating_percentage: {
19        $multiply: [
20          {
21            $divide: [
22              "$reviews.customer_reviews.rating",
23              5,
24            ],
25          }
26        ]
27      }
28    }
29  }
30 ]
```



## PIPELINE OUTPUT

Sample of 10 documents

OUTPUT OPTIONS

```
_id: ObjectId('645b422cc874e8c18f64e0a2')
Id: 21
ASIN: "0790747324"
title: "The Time Machine"
rating_percentage: 100
```

```
_id: ObjectId('645b422cc874e8c18f64e0a2')
Id: 21
ASIN: "0790747324"
title: "The Time Machine"
rating_percentage: 80
```

```
_id: ObjectId('645b422cc874e8c18f64e0a2')
Id: 21
ASIN: "0790747324"
title: "The Time Machine"
rating_percentage: 80
```

# Conclusion

## Journey so far:

- Explore various features offered by MongoDB and Cosmos DB
- Preprocess raw data into meaningful and useful JSON format
- Extract co-purchase relationships between products

## Next steps:

- Answer various questions by executing queries-
  - Identify products frequently purchased together
  - Detected patterns and trends in purchasing behavior
- Analyze the effectiveness of MongoDB in handling graph data

# References

- <https://www.mongodb.com/docs/manual/tutorial/>
- <https://learn.mongodb.com/>
- <https://techcrunch.com/2018/09/24/microsoft-updates-its-planet-scale-cosmos-db-database-service/>
- <https://www.mongodb.com/docs/manual/crud/>
- <https://www.mongodb.com/docs/manual/sharding/>
- <https://www.mongodb.com/basics/acid-transactions>
- <https://chat.openai.com/>
- <https://www.mongodb.com/docs/manual/core/aggregation-pipeline/>
- <https://www.mongodb.com/docs/compass/current/query/filter/>

# License

The materials for this presentation are licensed under CC BY-NC-SA  
(Attribution-NonCommercial-ShareAlike)

Creative Commons — Attribution-NonCommercial-ShareAlike 4.0 International — CC  
BY-NC-SA 4.0

Contact: [tufte@pdx.edu](mailto:tufte@pdx.edu)